Portfolio 1 - Utilitarianism vs. Deontology for Data Ethics: A Practical Comparison by Klaas-Jan Kisteman

Introduction

This short essay compares two classic ethical models—*Utilitarianism* and *Deontology*—and argues for a pragmatic preference in my own practice. Both aim to guide action, yet they evaluate moral rightness differently. In data science, where trade-offs between benefits, risks, and rules are frequent, understanding these differences is essential.

Core ideas

Utilitarianism judges actions by their consequences: the morally right act is the one that yields the greatest overall good (e.g., well-being) for the greatest number [1]. It is outcome-focused and naturally invites cost–benefit analysis, risk estimation, and impact assessment.

Deontology evaluates actions by duties, rules, or rights, independent of outcomes [2]. Certain actions (e.g., lying, coercion) can be wrong even if they would produce good consequences. This perspective often underpins strict constraints such as non-discrimination, respect for autonomy, or privacy duties.

Similarities and differences

Both models aim to systematize moral reasoning and provide action guidance. They differ primarily on what grounds moral status: *ends* (utilitarianism) versus *principles* (deontology). In practice, utilitarianism can flexibly balance competing interests and optimize outcomes, while deontology supplies bright lines that protect individuals against instrumental use. A known tension emerges in "hard cases": a utilitarian may endorse rule-breaking if the gains are large enough, whereas a deontologist resists this even under pressure.

My preferred model (practical stance)

For my work, I prefer a *utilitarian baseline with deontological side-constraints.* The baseline is practical: it aligns with how I already plan projects—estimating benefits (e.g., improved health outcomes), costs (e.g., false positives), and externalities (e.g., unfair burdens). It integrates well with tools I use (A/B tests, quantitative risk analysis, post-deployment monitoring).

However, I adopt *non-negotiable constraints* from deontology to prevent outcome-obsession from overriding essential rights: do not mislead stakeholders; avoid coercive dark patterns; never deploy models that violate consent or purpose limitation; and never justify discriminatory treatment by aggregate utility. These constraints narrow the search space so my optimization never comes at the expense of core duties.

Which suits data ethics best?

In data ethics, both models contribute. Utilitarian reasoning is indispensable for evaluating trade-offs at scale: model performance vs. fairness adjustments, data utility vs. privacy loss, or automation gains vs. displacement harms. It supports explicit documentation of impacts and proportionality analysis [3].

Yet deontological constraints are crucial for legitimacy and trust. Privacy, transparency, and non-discrimination are not merely variables; they are duties that structure permissible design choices. When privacy law or fairness norms apply, a duty-first filter prevents rationalizing harmful shortcuts.

Conclusion. The most workable approach for me is a hybrid: use utilitarianism for practical analysis and prioritization, bounded by deontological duties that protect rights and maintain integrity.

References (APA style)

References

[1] Wikipedia. (n.d.). *Utilitarisme*. In *Wikipedia, de vrije encyclopedie*. From `https://nl.wikipedia.org/wiki/Utilitarisme`

[2] Wikipedia. (n.d.). *Deontologie*. In *Wikipedia, de vrije encyclopedie*. From `https://nl.wikipedia.org/wiki/Deontologie`

[3] Mittelstadt, B. D., Floridi, L. (2016). *The ethics of big data: Current and foreseeable issues in biomedical contexts*. Science and Engineering Ethics, 22(2), 303–341. https://doi.org/10.1016/j.techfore.2016.10.013